Online Result Summary

Model: text_recognition

GPU(s): 2 x NVIDIA A100-SXM-80GB

Total Available GPU Memory: 158.6 GB

Constraint targets: Max p99 Latency : 10 ms

In 274 measurements across 30 configurations, **text_recognition_config_5** provides the best throughput: **9960 infer/sec**.

This is a **30% gain** over the default configuration (7671 infer/sec), under the given constraints on GPU(s) 2 x NVIDIA A100-SXM-80GB.

- **text_recognition_config_5**: 2 GPU instances with a max batch size of 32 on platform onnxruntime_onnx

Curves corresponding to the 3 best model configuration(s) out of a total of 30 are shown in the plots.

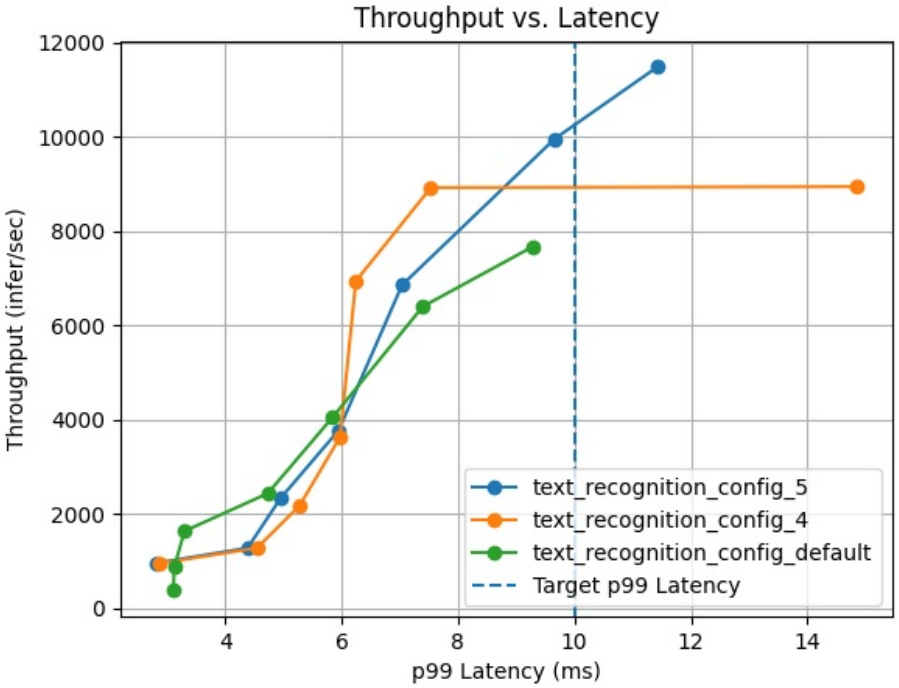

Throughput vs. Latency curves for 3 best configurations.

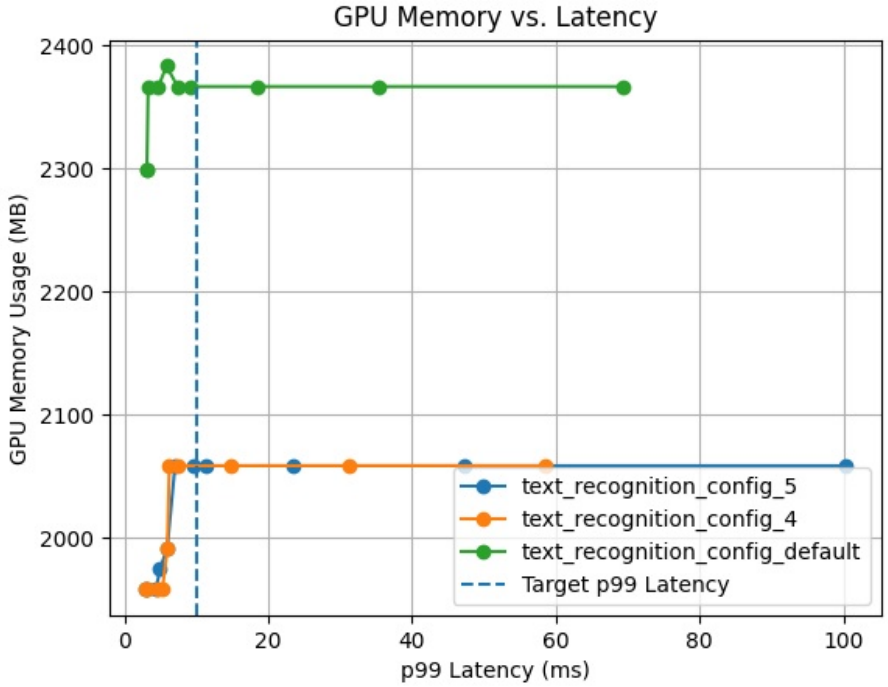

GPU Memory vs. Latency curves for 3 best configurations.

The following table summarizes each configuration at the measurement that optimizes the desired metrics under the given constraints.

Model Config Name	Max Batch Size	Dynamic Batching	Instance Count	p99 Latency (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
text_recognition_config_5	32	Enabled	2:GPU	9.658	9959.67	2058	82.6
text_recognition_config_4	16	Enabled	2:GPU	7.532	8925.97	2058	61.4
text_recognition_config_default	8	Enabled	4:GPU	9.279	7670.59	2366	72.0